

(Traditional) Object Detection

Daniil Osokin

Intel, IOTG/VMC/ICV

daniil.osokin@intel.com

Internet of Things Group

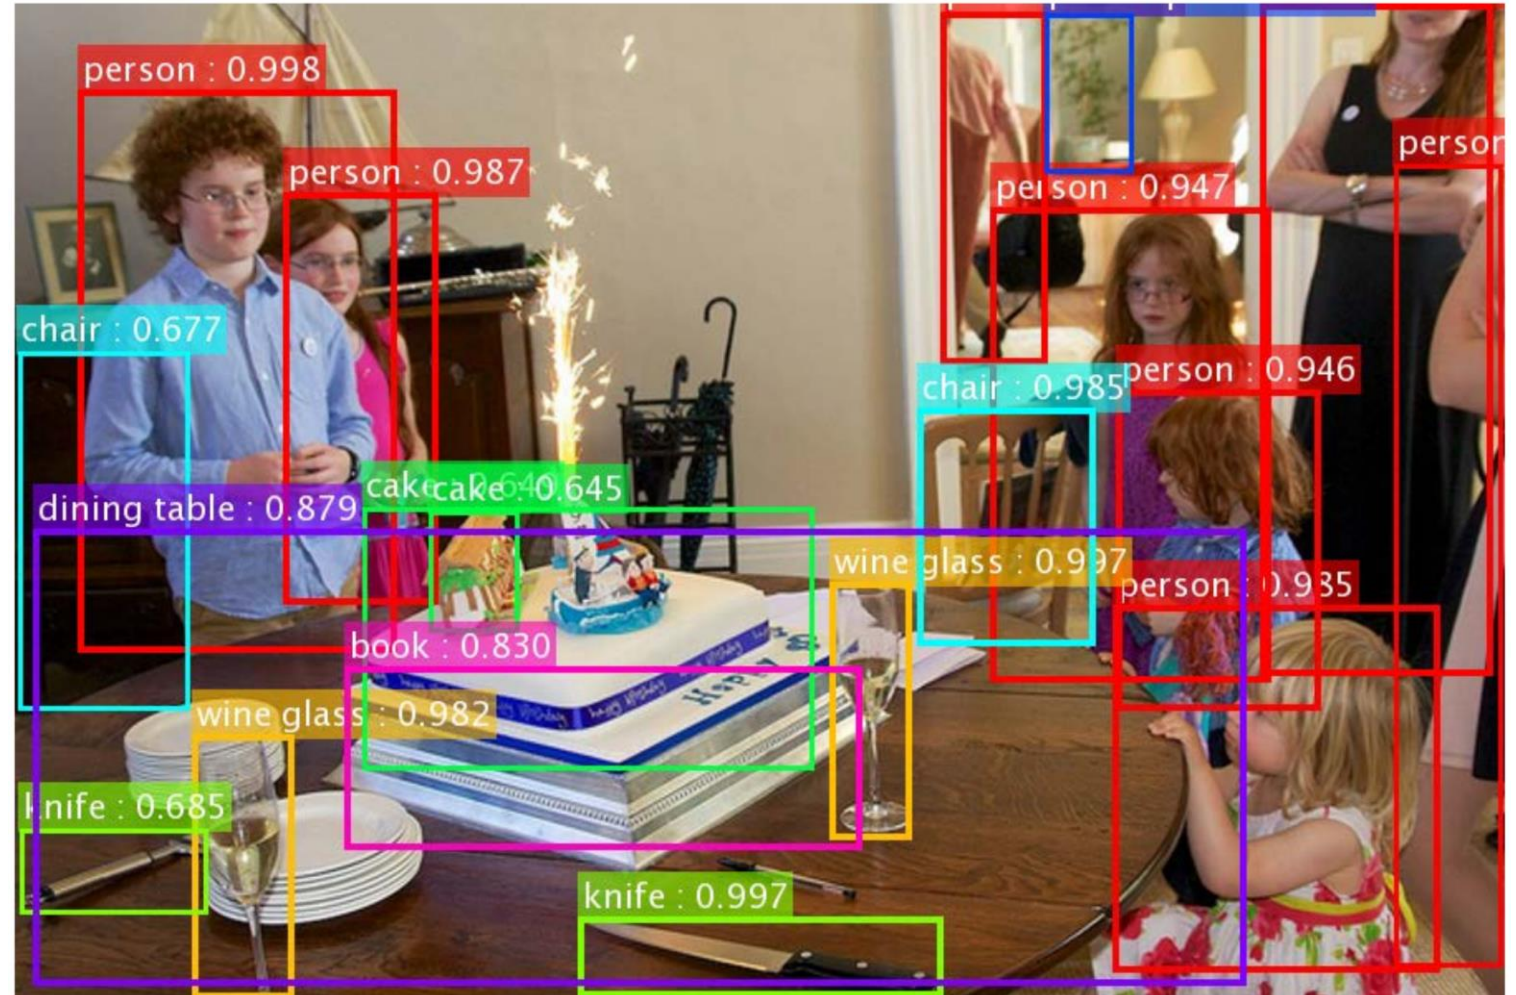
Object Detection

1. Localization:

- Make hypotheses about objects locations

2. Classification:

- Give label for found objects



Images credit: <https://chaosmail.github.io/deeplearning/2016/10/22/intro-to-deep-learning-for-computer-vision/>

Challenges

- Large variance in object appearance
 - Rotation (in-plane, out-of-plane), pose (rigid/non-rigid objects), shape, occlusion, environmental conditions, etc.
- Objects can have different scales on the same image
- Real-time speed for real-world applications

Classification Task

Train dataset: set of pairs $(\mathcal{X}_i, \mathcal{Y}_i)$, $i=1..N$ (training samples)

- \mathcal{X}_i - object, $\mathcal{X}_i = (x_i^1, x_i^2, \dots, x_i^n) \in \mathbb{R}^n$ – features vector (descriptor of object)
- $\mathcal{Y}_i \in \{0, 1\}$ – label for each object (class of object)

Test dataset: \mathcal{X}_i^{test} , *no* labels given

The classification task is to predict labels for each object in test dataset (do classification)



Decision Tree

Given train dataset $(\mathbf{x}_i, \mathbf{y}_i)$, $i=1..N$,
build a *tree* which reduces the
classification error:

$$\sum_{i=1}^N |label_{gt} - label_{predicted}|$$

$$\mathbf{x}_i = (x_i^1, x_i^2, \dots, x_i^n)$$

Decision tree (depth=1) examples:



Object Classification with Decision Tree

Decision tree (depth 1):
$$h_t(x) = \begin{cases} 1, & p^j x^j < p^j \theta^j \\ 0, & \text{otherwise} \end{cases}$$
 θ^j – threshold
 p^j – sign

Negative – object with label = 0

Positive – object with label $\neq 0$

Classifier predictions:

True Positive:

label_{gt} = 1
label_{predicted} = 1

False Positive:


label_{gt} = 0
label_{predicted} = 1

True Negative:

label_{gt} = 0
label_{predicted} = 0

False Negative:

label_{gt} = 1
label_{predicted} = 0



AdaBoost: Boosting with Decision Trees

1. Weighted train dataset: $(\mathbf{x}_i, \mathbf{y}_i)$, $i=1..N$, weights $w_i = \frac{1}{N}$ for each object

2. For $t = 1..T$

1. Train decision trees h_j , each makes split by x_i feature, compute error ε_j : $\varepsilon_j = \sum_{i=1}^m w_i * |h_j(x_i) - y_i|$

2. Select tree with the lowest error $\varepsilon_j = \varepsilon$

3. Update the weights of training samples: $w_i = w_i * \beta_t^{1-e_i}$, $\beta_t = \frac{\varepsilon}{1 - \varepsilon}$

$e_i = 0$ if \mathbf{x}_i classified correctly, else $e_i = 1$

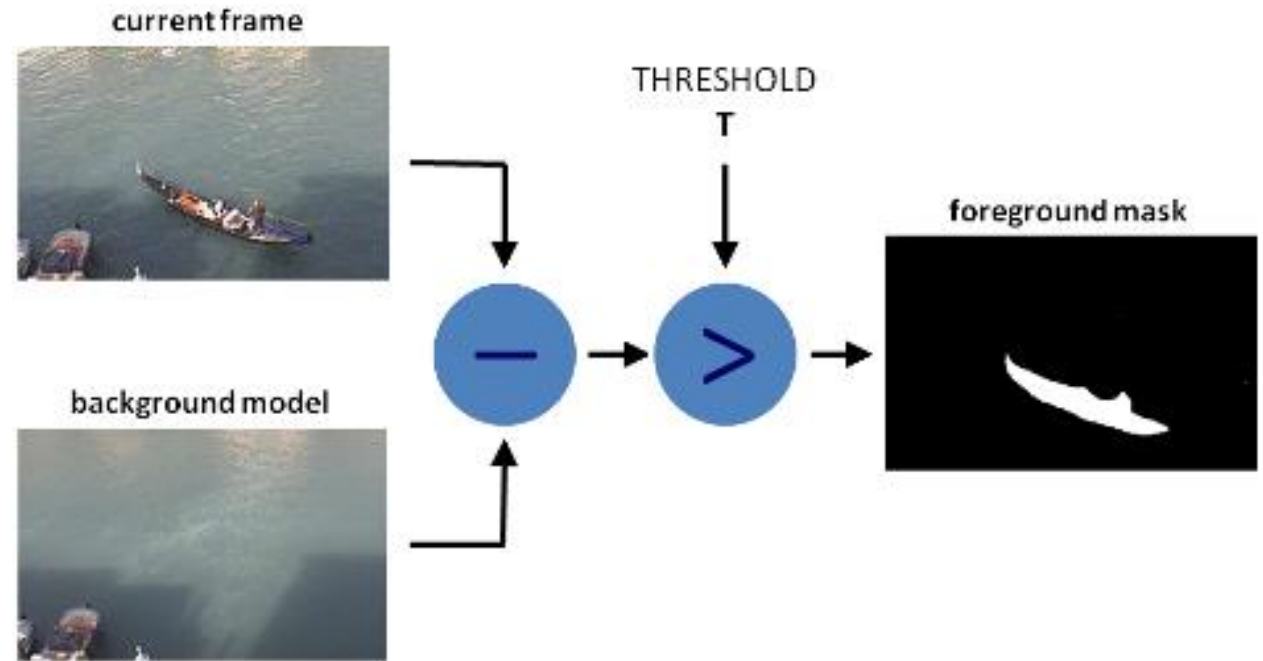
4. Normalize weights

3. Final (strong) classifier: $H(\mathbf{x}) = \begin{cases} 1, & \sum_{t=1}^T \alpha_t h_t(\mathbf{x}) \geq \frac{1}{2} \sum_{t=1}^T \alpha_t \\ 0, & \text{otherwise} \end{cases}$, $\alpha_t = \log \frac{1}{\beta_t}$

For pedestrian detection typical $T \sim 2000$

Localization: Background Subtraction

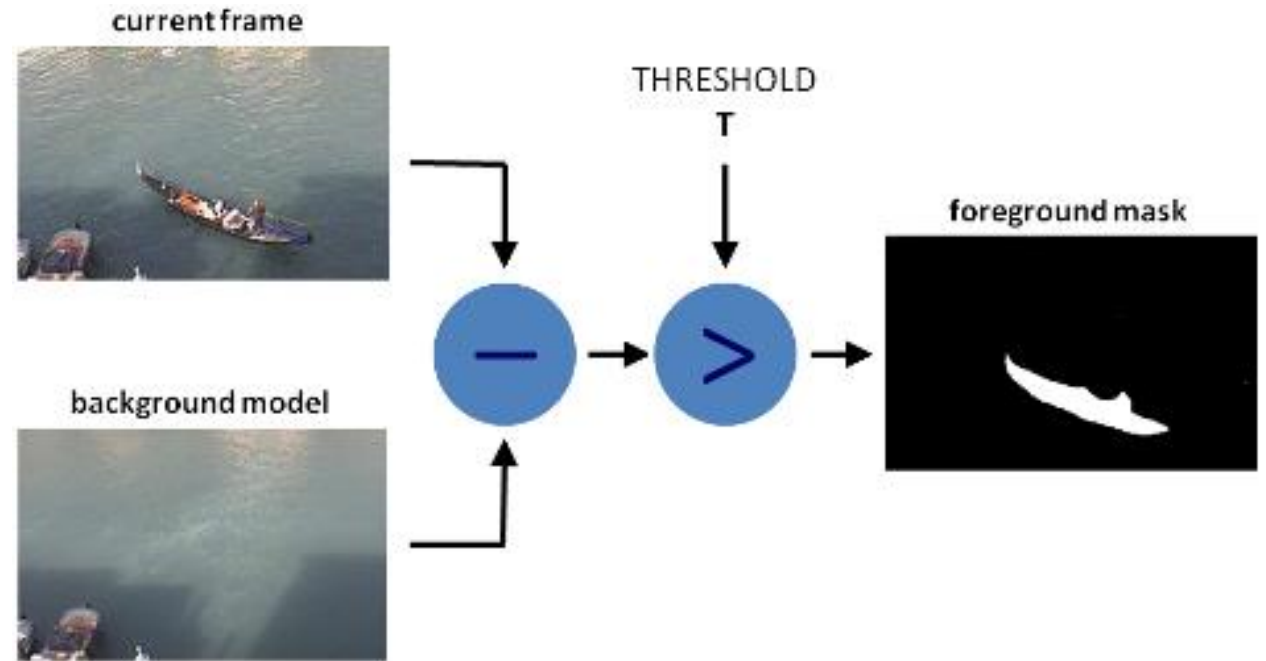
1. Given image of static background
2. Subtract from current frame the background, threshold to filter noise
3. Result mask contains *location hypotheses*



Images credit: <https://chaosmail.github.io/deeplearning/2016/10/22/intro-to-deep-learning-for-computer-vision/>

Localization: Background Subtraction

1. Given image of static background
2. Subtract from current frame the background, threshold to filter noise
3. Result mask contains *location hypotheses*



Q: When this does not work?


Images credit: <https://chaosmail.github.io/deeplearning/2016/10/22/intro-to-deep-learning-for-computer-vision/>

Localization: Sliding Window



1. Set fixed size window
2. Slide this window (left to right, top to bottom) over the image. Each window position is a *location hypothesis*
3. To deal with various scales, slide window over image pyramid



Face: 24x24 pixels



Person: 60x120 pixels



Images credit: <http://mi.eng.cam.ac.uk/~cipolla/lectures/4F12/Slides/old/4F12-detection.pdf>

Localization: Sliding Window

1. Set fixed size window
2. Slide this window (left to right, top to bottom) over the image. Each window position is a *location hypothesis*
3. To deal with various scales, slide window over image pyramid


Q: How many sliding windows for person in 1280x720 image



Face: 24x24 pixels



Person: 60x120 pixels



...



...




Images credit: <http://mi.eng.cam.ac.uk/~cipolla/lectures/4F12/Slides/old/4F12-detection.pdf>


Cascade Classifier

Some location hypotheses can be easily rejected, classified as negative, such as sky, road, uniform regions

No need to evaluate all classifiers in ensemble. Let's split it in *stages*.





Cascade Classifier: Example



Images credit: <http://mi.eng.cam.ac.uk/~cipolla/lectures/4F12/Slides/old/4F12-detection.pdf>

Cascade Classifier: Example



Face



188

hypotheses

Not face

Images credit: <http://mi.eng.cam.ac.uk/~cipolla/lectures/4F12/Slides/old/4F12-detection.pdf>

Cascade Classifier: Example





Not face

Face

48
hypotheses

Images credit: <http://mi.eng.cam.ac.uk/~cipolla/lectures/4F12/Slides/old/4F12-detection.pdf>

Cascade Classifier: Example



Face

18

Faces

Not face

Images credit: <http://mi.eng.cam.ac.uk/~cipolla/lectures/4F12/Slides/old/4F12-detection.pdf>

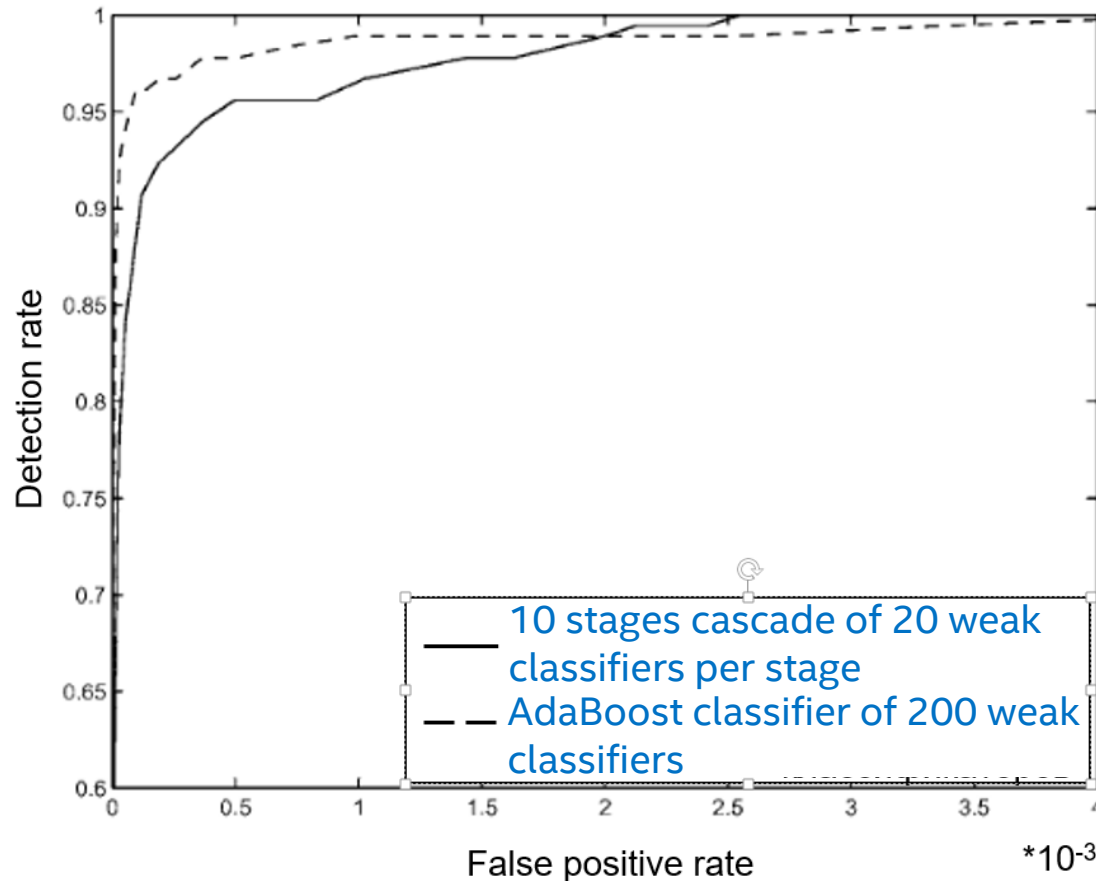
Comparison of AdaBoost with Cascade Classifier (Viola & Jones, 2001)



Paul Viola,
Microsoft
research



Michael Jones,
Merl



Cascade boosts performance in ~ 10 times with comparable accuracy

384x288 pixels input

Intel Pentium III 700 Mhz

@ 15 fps



Images credit: <http://www.vision.caltech.edu/html-files/EE148-2005-Spring/pprs/viola04ijcv.pdf>

Can you find a face in the coffee beans?



Images credit: <http://mi.eng.cam.ac.uk/~cipolla/lectures/4F12/Slides/old/4F12-detection.pdf>

Features

How to form feature vector?

pixels values



$$\longrightarrow \mathcal{X} = (50, 54, 52, 54, \dots, 186)$$

Feature vector, 512x1 пикселей

gradient values



$[-1 \ 0 \ 1]$

$$\longrightarrow \mathcal{X} = (0, 0, 38, 0, \dots, 0)$$

Feature vector, 512x1 пикселей

Features

How to form feature vector?

pixels values



$$\longrightarrow \mathcal{X} = (50, 54, 52, 54, \dots, 186)$$

Feature vector, 512x1 пикселей

gradient values



$[-1 \ 0 \ 1]$

$$\longrightarrow \mathcal{X} = (0, 0, 38, 0, \dots, 0)$$

Feature vector, 512x1 пикселей

Which feature better?


Features: Haar-like

Features should be fast to compute at any scale




1. Edge features



2. Line features





3. Center-surround features




Images credit: <https://github.com/itseez-academy/itseez-summer-school-2015-lectures/blob/master/slides/Day%203%20--%20Alexander%20Bovyrin%20--%20Object%20Detection.pdf>

Integral Image for Fast Computation

$$ii(x, y) = \sum_{x' \leq x, y' \leq y} i(x', y'),$$





Input image




Integral image

Sum inside:

 = $(4 + 0) - (3 + 0) = 1$


 = $(12 + 0) - (0 + 8) = 4$

 = $(12 + 1) - (4 + 3) = 6$

Features: Histogram of Oriented Gradients (HoG)

1. Compute horizontal, vertical gradients for each pixel inside the image
2. Compute magnitude, orientation of gradients
3. Divide image (size 64x128) by blocks 16x16 pixels: 2x2 cells by 8x8 pixels (7 * 15 blocks in total)





4. For each block compute histogram of gradient orientation



1. Kernels of gradients
2. Magnitude, orientation (angle)


5. Normalize blocks

6. Concatenate blocks histograms into feature vector



Images credit: <http://crv.ucf.edu/courses/CAP5415/Fall2012/Lecture-6a-Hog.pdf>

Features: Histogram of Oriented Gradients (HoG)



Magnitude of gradient in specific direction is visualized in green

Images credit: <http://crcv.ucf.edu/courses/CAP5415/Fall2012/Lecture-6a-Hog.pdf>

Summary

Classification:

- Decision Tree
- AdaBoost
- Cascade Classifier

Localization:

- Sliding window
- Image pyramid

Features:

- Haar-like
- Histogram of Oriented Gradients (HoG)

Q&A

daniil.osokin@intel.com