

Neural Network Optimization for Efficient Inference

Alexander Suslov


Neural Network Applications

Self-Driving Car



This image is in the public domain

Robots



This image is in the public domain

Image processing




This image is in the public domain

Machine Translation




This image is in the public domain

Medical imaging



This image is in the public domain

3D scanning



This image is in the public domain

Where does Inference of Neural Networks Compute?

Standalone




Client-Server



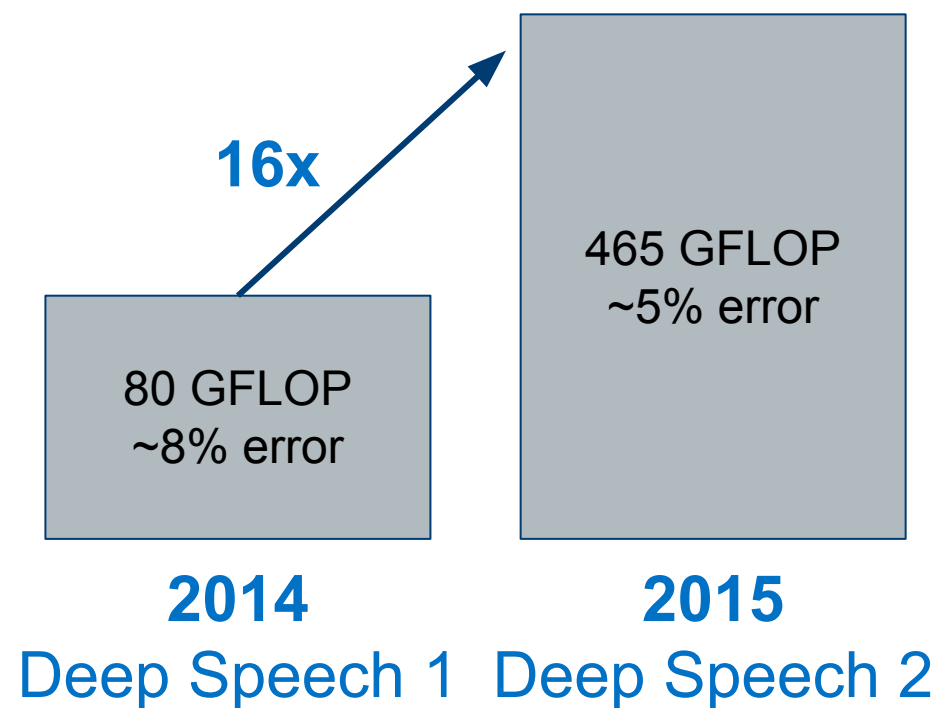
All images are in the public domain

Models are Getting Larger

Image Classification




Speech Recognition

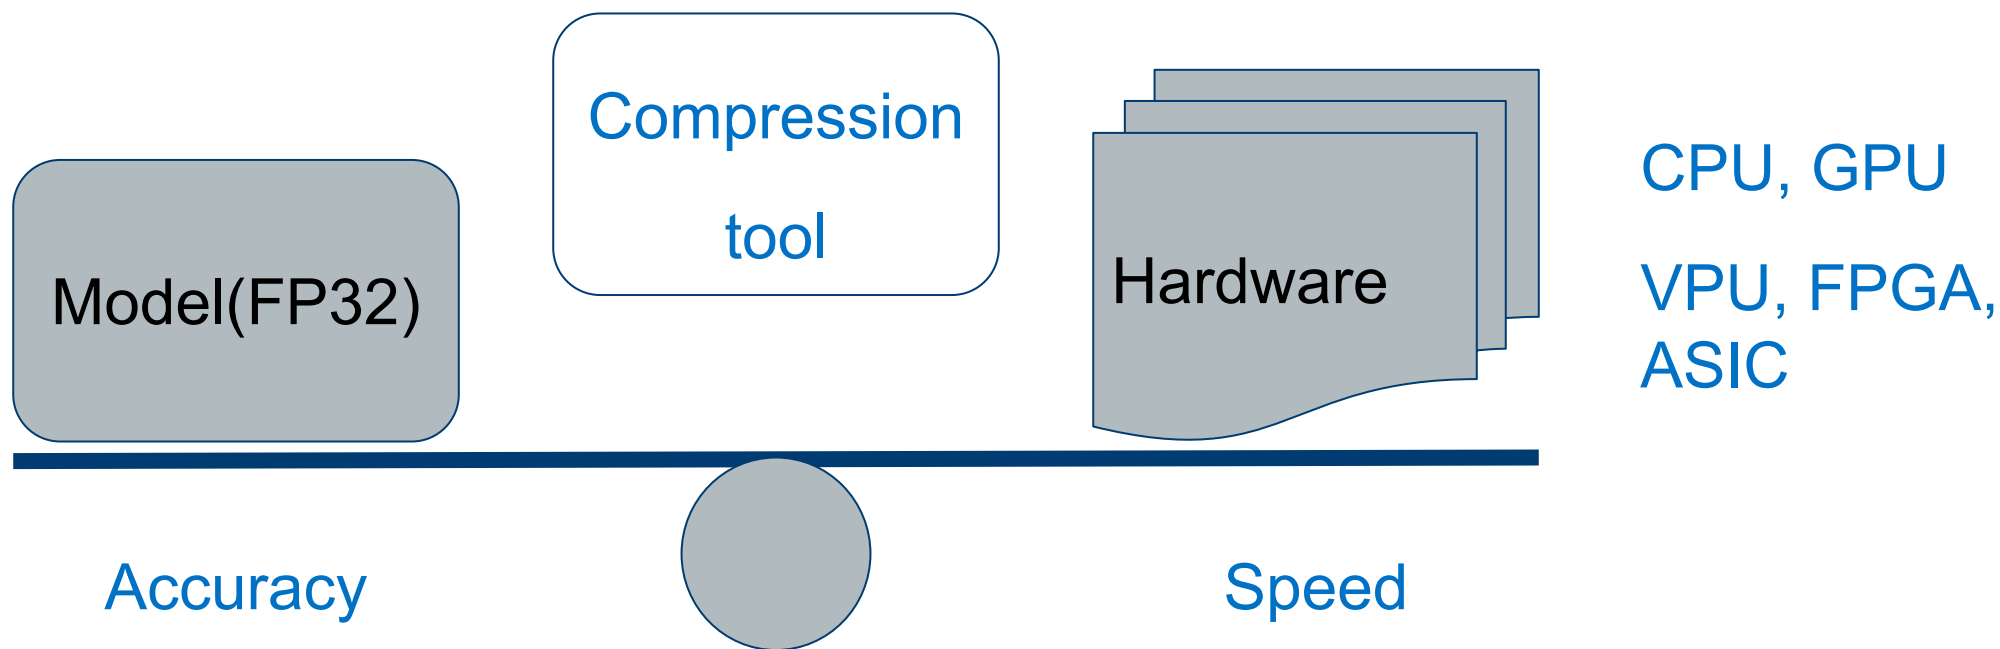


The First Challenge: Model Size

- Hard to distribute large models through over-the-air update
- The first run is slow due to loading weights.



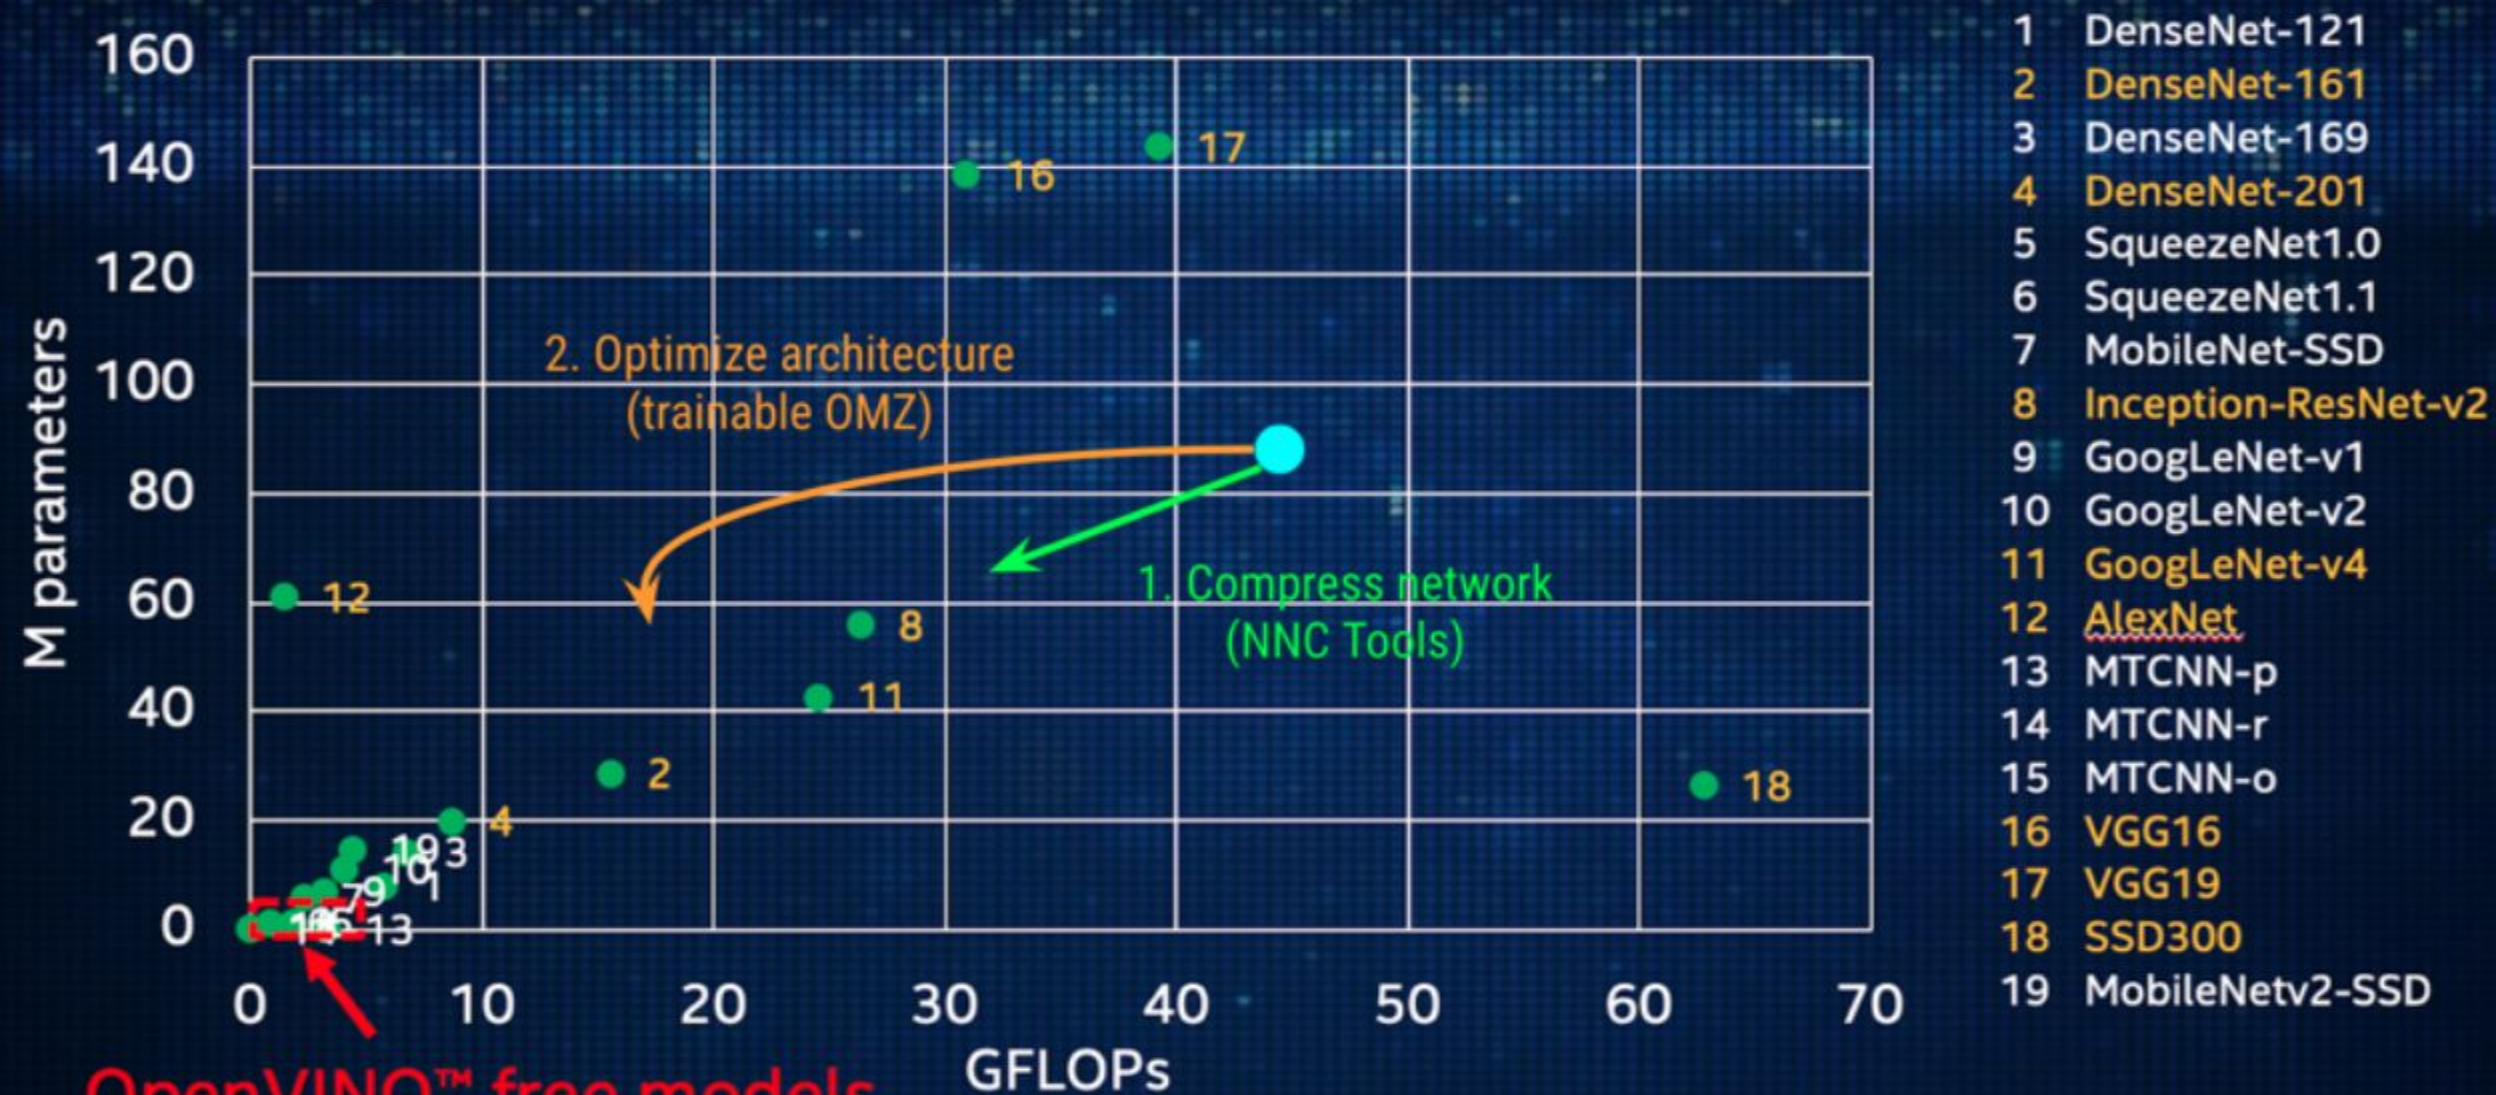
The Second Challenge: Speed



Trade off between accuracy and performance

Algorithms for Efficient Inference


1. Pruning
2. Weight Sharing
3. Quantization
4. Binary / Ternary Net
5. Distillation
6. Low Rank Approximation
7. Winograd Transformation



Algorithms for Efficient Inference

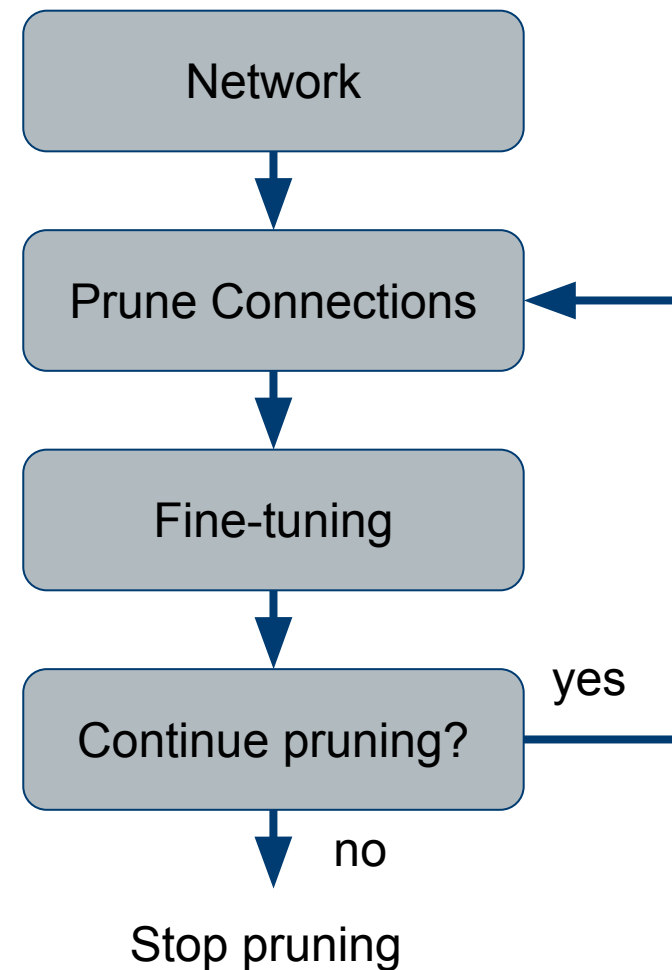
1. Pruning
2. Weight Sharing
3. Quantization
4. Binary / Ternary Net
5. Distillation
6. Low Rank Approximation
7. Winograd Transformation

Pruning Neural Networks





Pruning Neural Networks

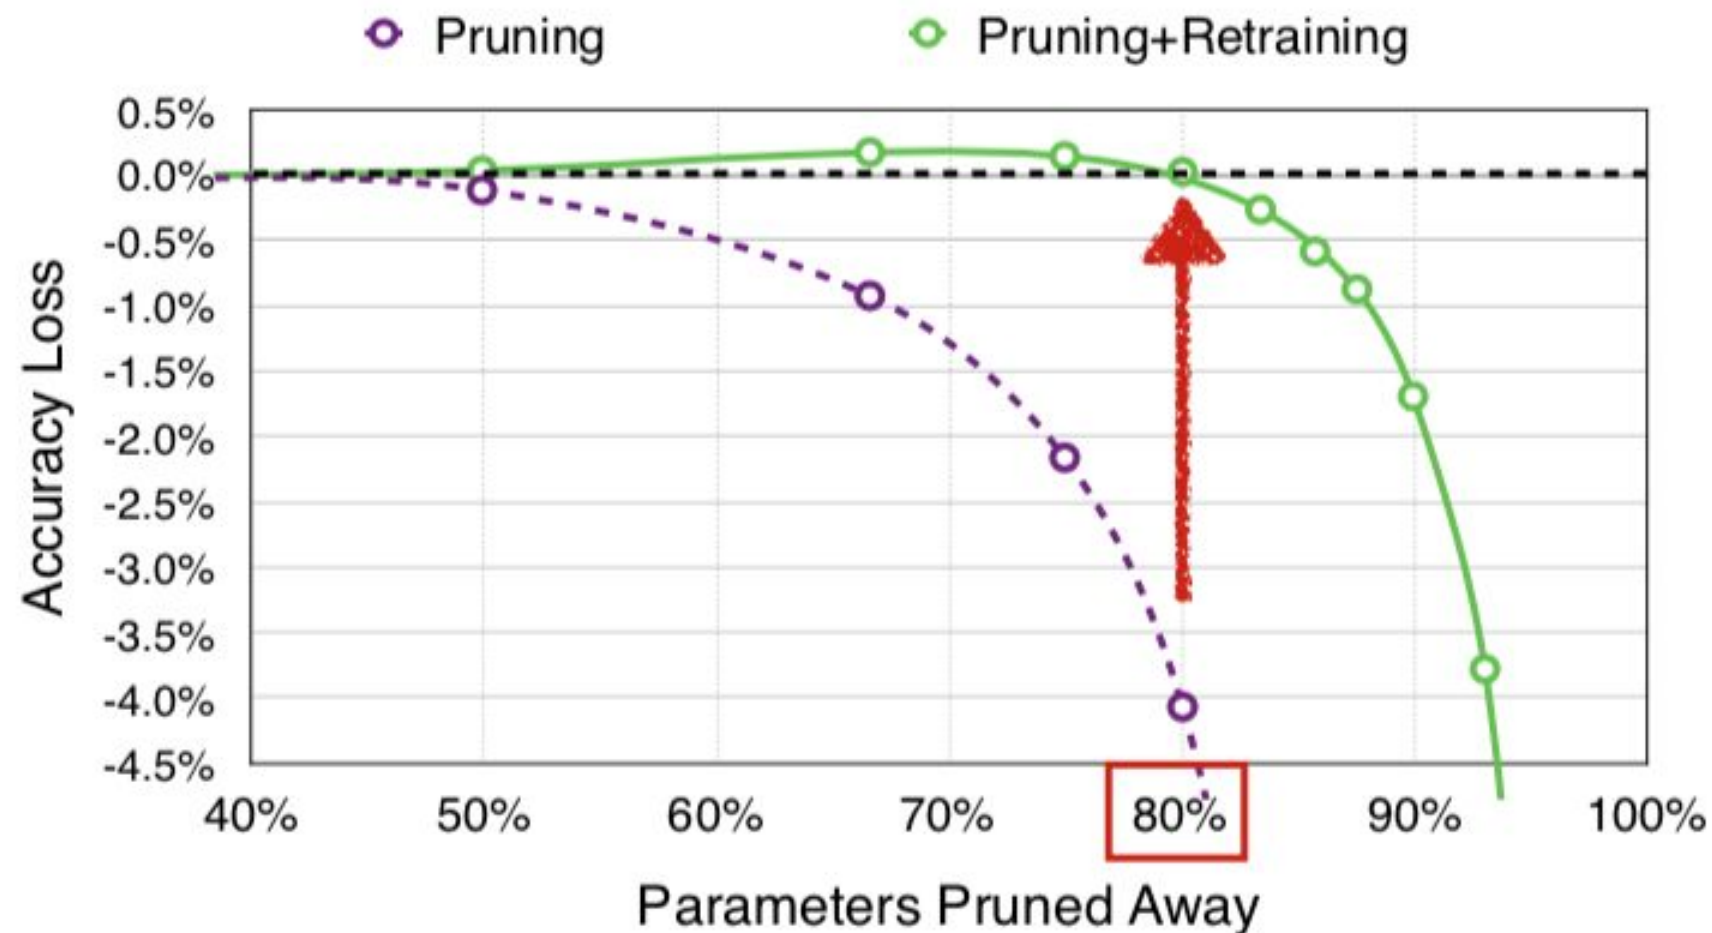

- Criteria for pruning
 - Connections with low weights
 - Neurons or filters with low impact
- External limitations
 - Spatial structure of pruned weights





Pruning Neural Networks



Retrain to Recover Accuracy

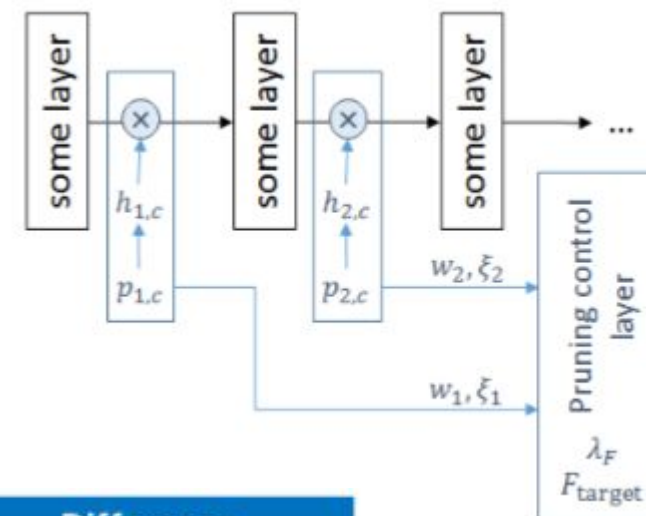


Iteratively Retrain to Recover Accuracy



Channel Pruning

- Drop some channels by Dropout layers with learnable dropout probabilities.



Task	Original network	Pruned network	Difference
VGG-16 ILSVRC2012 classification (1000 classes)	15.47 GFLOP 68.4% top-1 acc. 88.4% top-5 acc.	3.87 GFLOP 67.4% top-1 acc. 88.0% top-5 acc.	-75% (FLOP) -1.0% top-1 acc. -0.4% top-5 acc.
MobileNet + SSD Vehicle detection (1 class)	2.26 GFLOP 25 ms@CPU 87.7% AP	1.33 GFLOP 19 ms@CPU 86.2% AP	-41% (FLOP) -1.5% AP
AlexNet ILSVRC2012 classification (1000 classes)	724 MFLOP 56.8% top-1 acc. 79.9% top-1 acc.	411 MFLOP 55.8% top-1 acc. 79.1% top-5 acc.	-43% (FLOP) -1.0% top-1 acc. -0.8% top-5 acc.

Sparsification

- The sparsification algorithm based on probabilistic approach and loss regularization.

Inception v3

Baseline: 77.46% top1 acc.

Sparsity Rate(SR) 50%: 77.25% top1 acc.

SR 92%: 76.6% top1 acc.

MobileNet v2

Baseline: 71.6% top1 acc.

SR 50%: 71.2% top1 acc.

SR 78%: 69.98% top1 acc.

Ordinary convolution

$$output = conv(x, w)$$

Sparsifying weights we reparametrize weights as:

$$\bar{w} = w * z$$

Where:

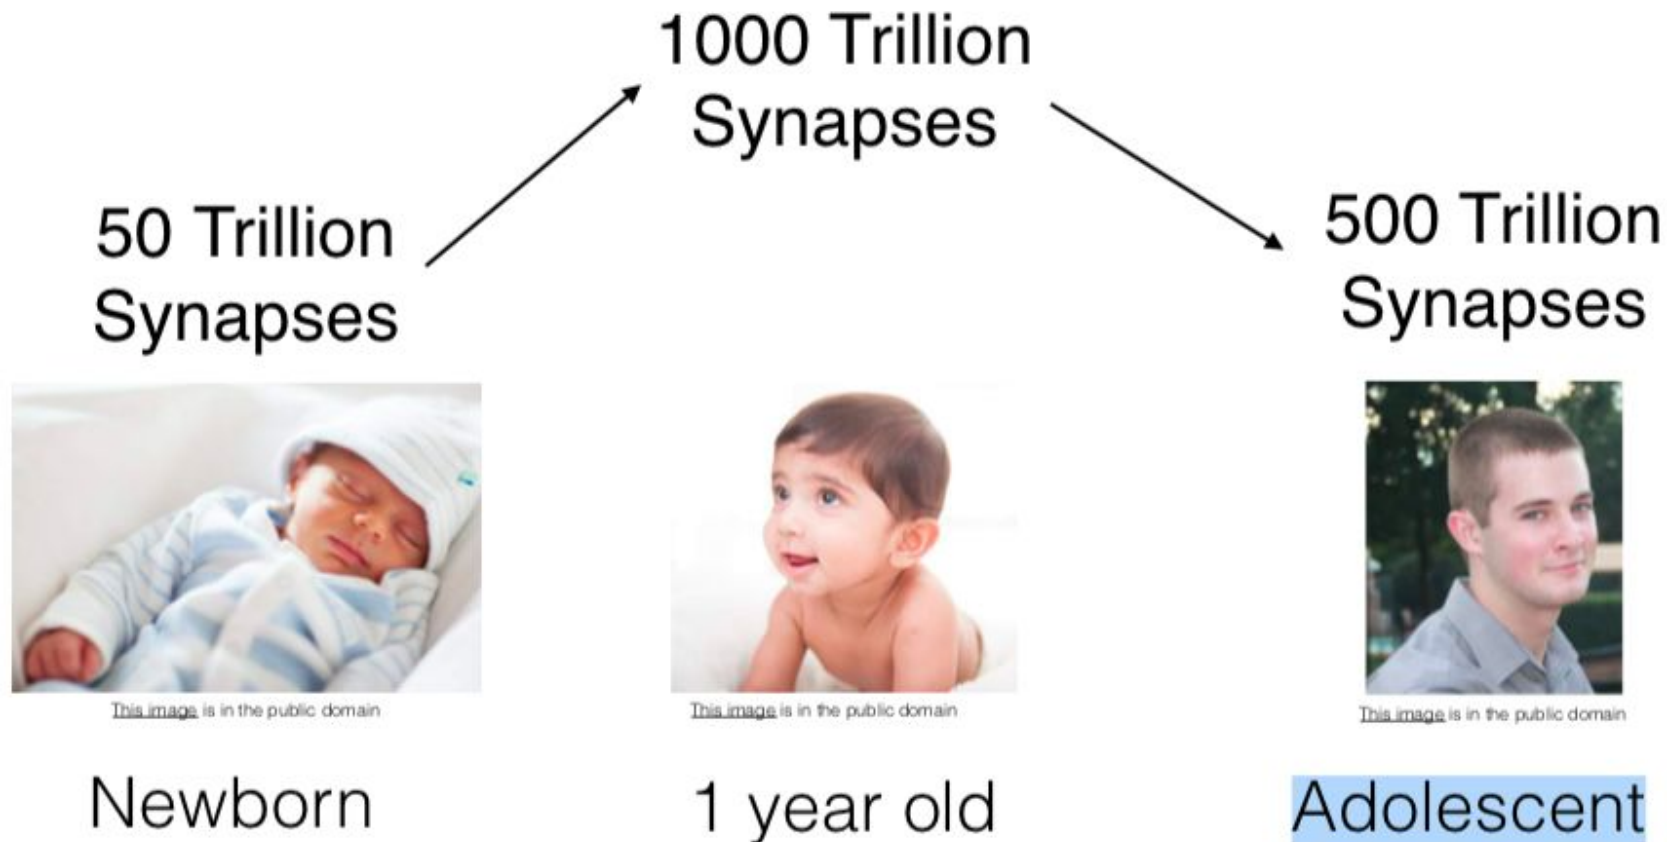
w – weights, $w \in R$

z – binary mask, $z \in [0, 1]$

We train these masks using modified loss

$$Loss = Loss_{task} + \alpha \left(\sum_l^{Layers} ||z|| - target \right)^2$$

Pruning Happens in Human Brain




Christopher A Walsh. Peter Huttenlocher (1931-2013). Nature, 502(7470):172–172, 2013.

Slide credits by Song Han

Algorithms for Efficient Inference

1. Pruning
2. **Weight Sharing**
3. Quantization
4. Binary / Ternary Net
5. Distillation
6. Low Rank Approximation
7. Winograd Transformation

Weight sharing




Pruning + Weight sharing + Huffman Encoding

Network	Original Size	Compressed Size	Compression Ratio	Original Accuracy	Compressed Accuracy
LeNet-300	1070KB → 27KB		40x	98.36%	→ 98.42%
LeNet-5	1720KB → 44KB		39x	99.20%	→ 99.26%
AlexNet	240MB → 6.9MB		35x	80.27%	→ 80.30%
VGGNet	550MB → 11.3MB		49x	88.68%	→ 89.09%
GoogleNet	28MB → 2.8MB		10x	88.90%	→ 88.92%
ResNet-18	44.6MB → 4.0MB		11x	89.24%	→ 89.28%

Algorithms for Efficient Inference


1. Pruning
2. Weight Sharing
3. Quantization
4. Binary / Ternary Net
5. Distillation
6. Low Rank Approximation
7. Winograd Transformation

Quantization



Quantization

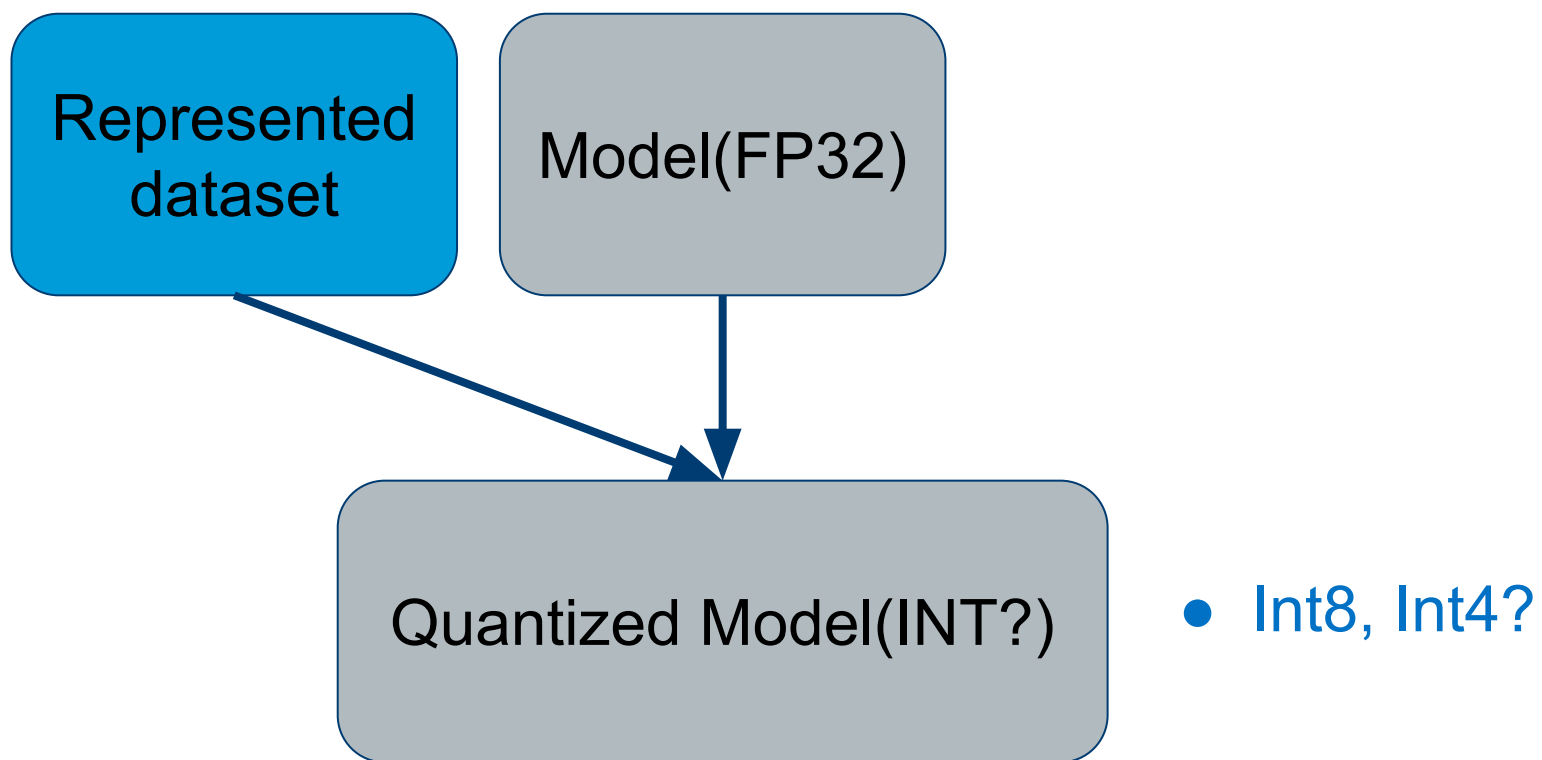
- Level 1: Post-Training Quantization w/o dataset



- Int8
- BatchNorms

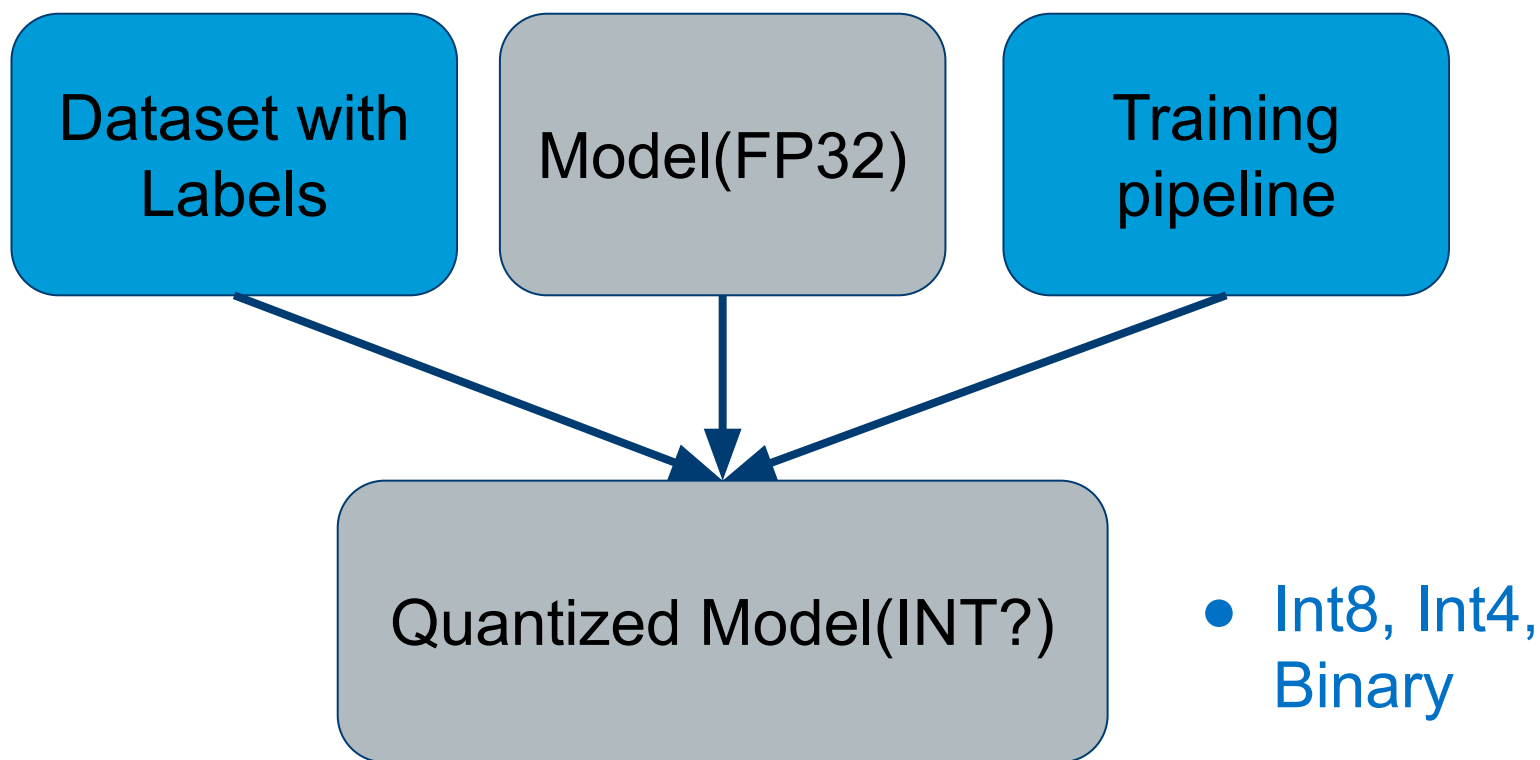
Quantization

- Level 2: Post-Training Quantization w dataset



Quantization

- Level 3: Quantization via training



Uniform Affine Quantization

- Quantization:

$$x_{int} = \text{round}\left(\frac{x}{\Delta}\right) + z$$

$$x_Q = \text{clamp}(0, N_{levels} - 1, x_{int})$$

- Δ specifies the step size of the quantizer and floating point zero maps to zero-point.
- z - zero-point.
- $N_{levels} = 256$ for 8-bits of precision

- De-quantization:

$$x_{float} = (x_Q - z)\Delta$$

Uniform Affine Quantization

- 2D convolution between a weight and an activation:

$$y(k, l, n) = \Delta_w \Delta_x \text{conv}(w_Q(k, l, m; n) - z_w, x_Q(k, l, m) - z_x)$$

$$y(k, l, n) = \text{conv}(w_Q(k, l, m; n), x_Q(k, l, m)) - z_w \sum_{k=0}^{K-1} \sum_{l=0}^{K-1} \sum_{m=0}^{N-1} x_Q(k, l, m) \\ - z_x \sum_{k=0}^{K-1} \sum_{l=0}^{K-1} \sum_{m=0}^{N-1} w_Q(k, l, m; n) + z_x z_w$$

Uniform Symmetric Quantization

- Quantization, zero-point = 0

- Activations:

$$x_{int} = \text{round}\left(\frac{x}{\Delta}\right)$$

$$x_Q = \text{clamp}(-N_{levels}/2, N_{levels}/2 - 1, x_{int}) \quad \text{if signed}$$

$$x_Q = \text{clamp}(0, N_{levels} - 1, x_{int}) \quad \text{if un-signed}$$

- Weights

$$x_Q = \text{clamp}(-(N_{levels}/2 - 1), N_{levels}/2 - 1, x_{int}) \quad \text{if signed}$$

$$x_Q = \text{clamp}(0, N_{levels} - 2, x_{int}) \quad \text{if un-signed}$$

Uniform Symmetric Quantization

- MKL DNN Int8 Workflow:

$$X_{s32} = W_{s8} \times \alpha_{u8} + b_{s32} \approx Q_{\alpha} Q_w X_{f32}$$

$$\text{where } X_{f32} = W_{f32} \times \alpha_{f32} + b_{f32}$$

$Q_{\alpha} = \frac{255}{R_{\alpha}}$ is the quantization factor for activations with non-negative values.

$Q_w = \frac{127}{R_w}$ is the quantization factor for weights.

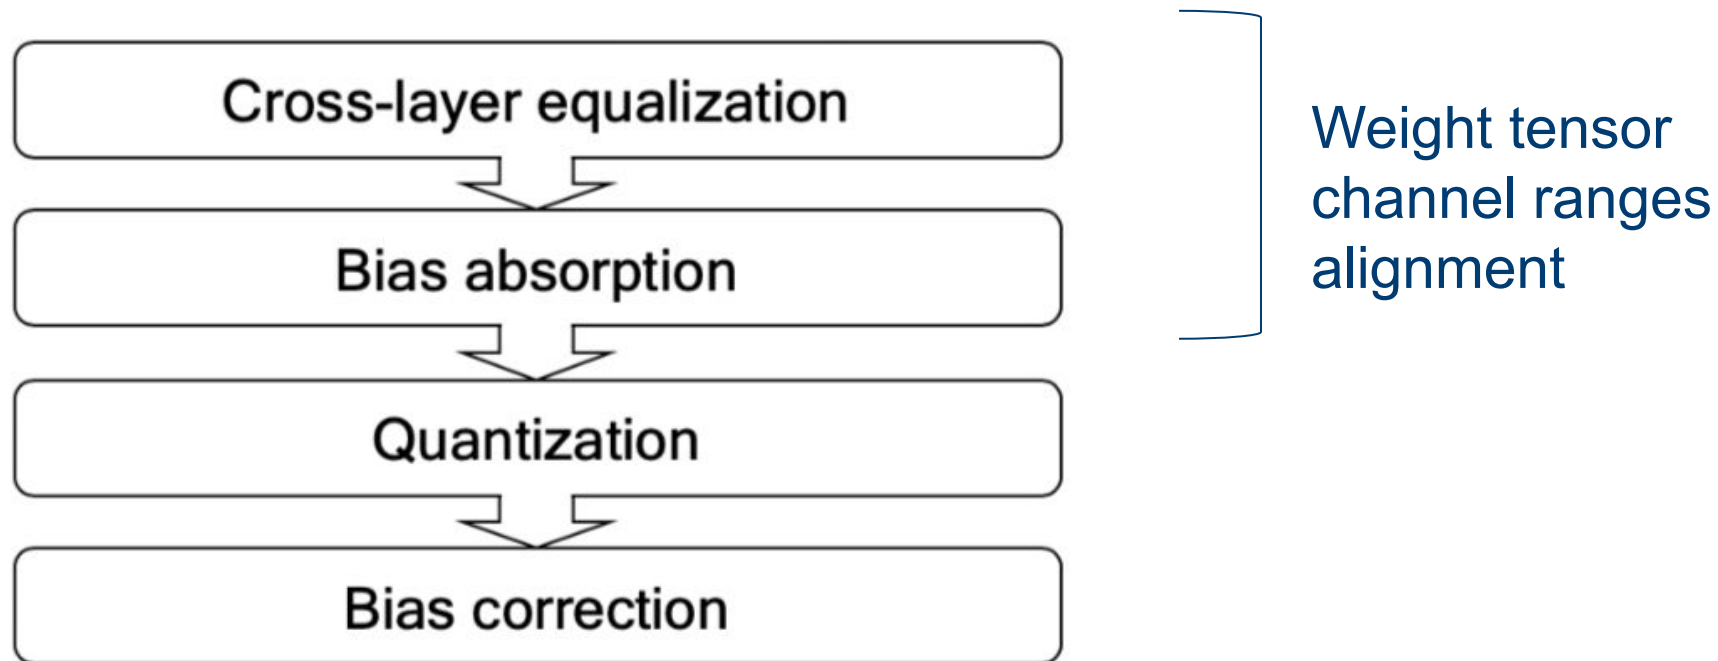
$$\alpha_{u8} = \lceil Q_{\alpha} \alpha_{f32} \rceil \in [0, 255]$$

$$W_{s8} = \lceil Q_w W_{f32} \rceil \in [-127, 127]$$

$$b_{s32} = \lceil Q_{\alpha} Q_w b_{f32} \rceil \in [-2^{31}, 2^{31} - 1]$$

Data-Free Quantization (Level 1)

- The basic idea is to use BatchNorm statistics to estimate the range of activations.
- Algorithm Flow:



Cross-Layer Equalization

Given two layers, $\mathbf{x}_1 = f(\mathbf{W}_1\mathbf{x}_0 + \mathbf{b}_1)$ and $\mathbf{x}_2 = f(\mathbf{W}_2\mathbf{x}_1 + \mathbf{b}_2)$ through scaling invariance we have that:

$$\begin{aligned}\mathbf{x}_2 &= f(\mathbf{W}_2 f(\mathbf{W}_1\mathbf{x}_0 + \mathbf{b}_1) + \mathbf{b}_2) \\ &= f(\mathbf{W}_2 \mathbf{S} \hat{f}(\mathbf{S}^{-1}\mathbf{W}_1\mathbf{x}_0 + \mathbf{S}^{-1}\mathbf{b}_1) + \mathbf{b}_2) \\ &= f(\widehat{\mathbf{W}}_2 \hat{f}(\widehat{\mathbf{W}}_1\mathbf{x}_0 + \widehat{\mathbf{b}}_1) + \mathbf{b}_2)\end{aligned}$$

where:

$$s_i = \frac{1}{r_i^{(1)}} \sqrt{r_i^{(1)} r_i^{(2)}} \quad ; \quad r_i^{(1)} = \underline{2} \cdot \max_j |\widehat{\mathbf{W}}_{ij}^{(1)}|$$

Quantization

- Uniform Affine Quantization is used for weights and activations
 - $(\min(w), \max(w))$ - range for weights
 - $(\text{mean} - 3 * \text{var}, \text{mean} + 3 * \text{var})$ - range for activations

Quantization bias correction

- \mathbf{x} -input
- \mathbf{W} - weights FP32
- $\tilde{\mathbf{W}}$ - quantized weights

$$\tilde{\mathbf{y}} = \mathbf{y} + \boldsymbol{\epsilon}\mathbf{x}, \quad \boldsymbol{\epsilon} = \tilde{\mathbf{W}} - \mathbf{W}$$

$$\begin{aligned}\mathbb{E}[\mathbf{y}] &= \mathbb{E}[\mathbf{y}] + \mathbb{E}[\boldsymbol{\epsilon}\mathbf{x}] - \mathbb{E}[\boldsymbol{\epsilon}\mathbf{x}] \\ &= \mathbb{E}[\tilde{\mathbf{y}}] - \mathbb{E}[\boldsymbol{\epsilon}\mathbf{x}]\end{aligned}$$

$$\begin{aligned}\mathbb{E}[\mathbf{x}_c] &= \mathbb{E}[\text{ReLU}(\mathbf{x}_c^{pre})] \\ &= \gamma_c \mathcal{N}\left(\frac{-\beta_c}{\gamma_c}\right) + \beta_c \left[1 - \Phi\left(\frac{-\beta_c}{\gamma_c}\right)\right]\end{aligned}$$

Quantization with represented dataset (Level 2)


- Run N examples through the FP32 model and collect for each layer the per-channel pre-activation statistics:
 - moving average of the minimum and maximum values across batches to determine the quantizer parameters for activations.
 - mean activation for quantization bias correction

Quantization via training (Level 3)

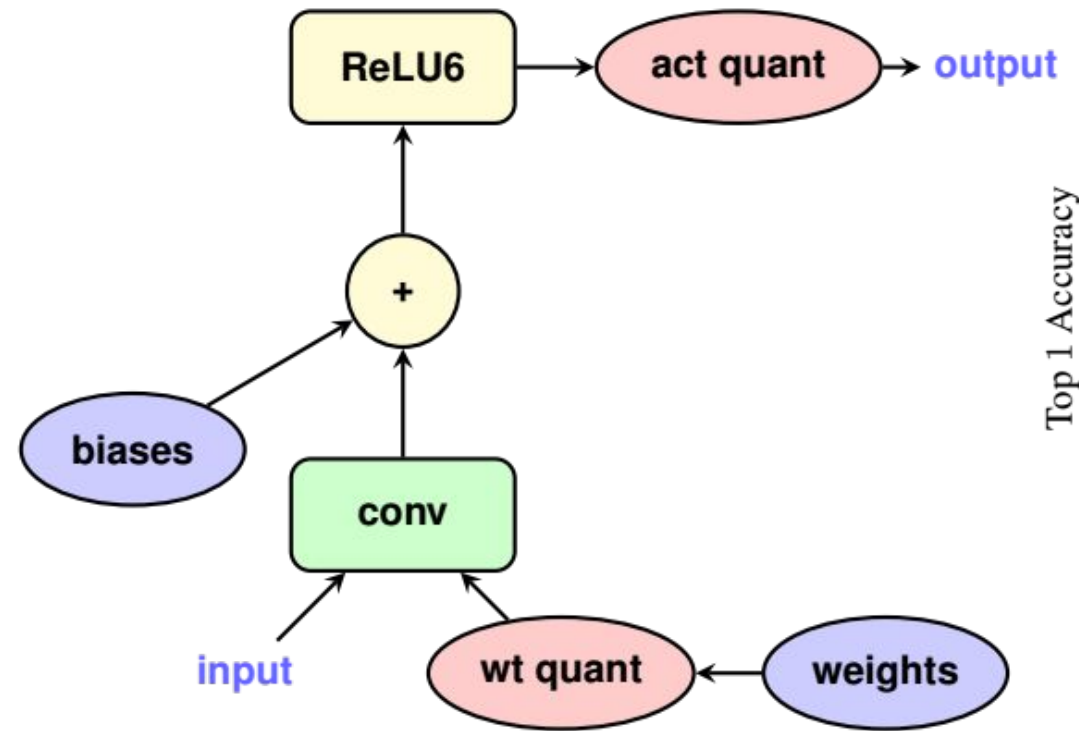
- Simulated quantization layer (Fake Quantization Layer)

$$\begin{aligned} \text{clamp}(r; a, b) &:= \min(\max(x, a), b) \\ s(a, b, n) &:= \frac{b - a}{n - 1} \\ q(r; a, b, n) &:= \left\lfloor \frac{\text{clamp}(r; a, b) - a}{s(a, b, n)} \right\rfloor s(a, b, n) + a, \end{aligned} \tag{12}$$

Quantization via training (Level 3)



(a) Integer-arithmetic-only inference



(b) Training with simulated quantization

Top 1 Accuracy

Quantization via training (Level 3)

Algorithm:

1. Create a training graph of the floating-point model.
2. Insert fake quantization operations in locations where tensors will be downcasted to fewer bits during inference.
3. Train in simulated quantized mode until convergence.
4. Create and optimize the inference graph for running in a low bit inference engine.
5. Run inference using the quantized inference graph.

Summary: Quantization approaches


	~D	~BP	~AC	MobileNetV2		MobileNetV1		ResNet18		
				FP32	INT8	FP32	INT8	FP32	INT8	INT6
DFQ (ours)	✓	✓	✓	71.7%	71.2%	70.8%	70.5%	69.7%	69.7%	66.3%
Per-layer [18]	✓	✓	✓	71.9%	0.1%	70.9%	0.1%	69.7%	69.2%*	63.8%*
Per-channel [18]	✓	✓	✓	71.9%	69.7%	70.9%	70.3%	69.7%	69.6%*	67.5%*
QT [16] ^	✗	✗	✓	71.9%	70.9%	70.9%	70.0%	-	70.3% [†]	67.3% [†]
SR+DR [†]	✗	✗	✓	-	-	-	71.3%	-	68.2%	59.3%
QMN [31]	✗	✗	✗	-	-	70.8%	68.0%	-	-	-
RQ [21]	✗	✗	✗	-	-	-	70.4%	-	69.9%	68.6%

Table 5. Top1 ImageNet validation results for different models and quantization approaches. The top half compares level 1 approaches (~D: data free, ~BP: backpropagation-free, ~AC: Architecture change free) whereas in the second half we also compare to higher level approaches in literature. Results with * indicates our own implementation since results are not provided, ^ results provided by [18] and [†] results from table 2 in [21].


Algorithms for Efficient Inference

1. Pruning
2. Weight Sharing
3. Quantization
4. Binary / Ternary Net
5. Distillation
6. Low Rank Approximation
7. Winograd Transformation

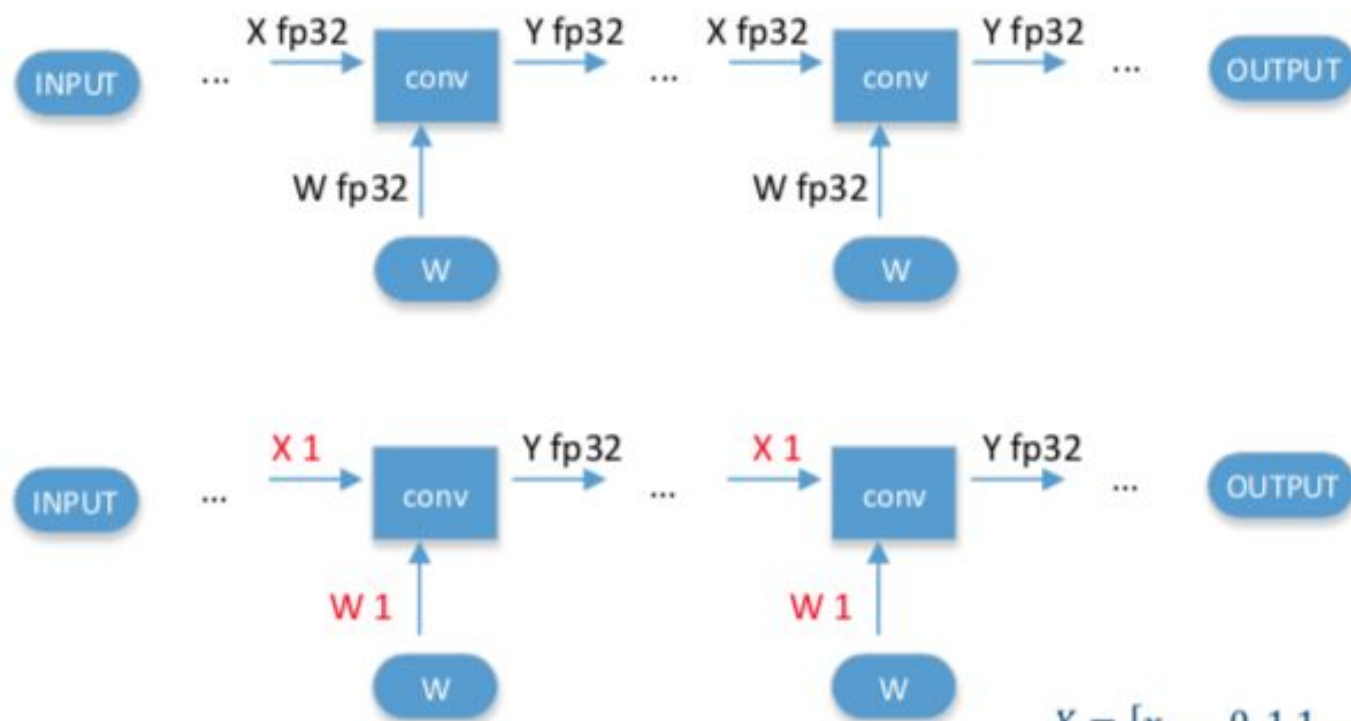
Binary / Ternary Net



\Rightarrow



Binary Net



1 value inference in conv is dot product

$$y_1 = w_1 * x_1 + w_2 * x_2 + \dots + w_{32} * x_{32}$$

32 * and 32+ operations

$$y_1 = \{-1, +1\} * \{-1, +1\} + w_2 * x_2 + \dots$$

$$x_i * w_i = \begin{matrix} x \backslash w & -1 & +1 \\ -1 & +1 & -1 \\ +1 & -1 & +1 \end{matrix} \quad \{-1, +1\},$$

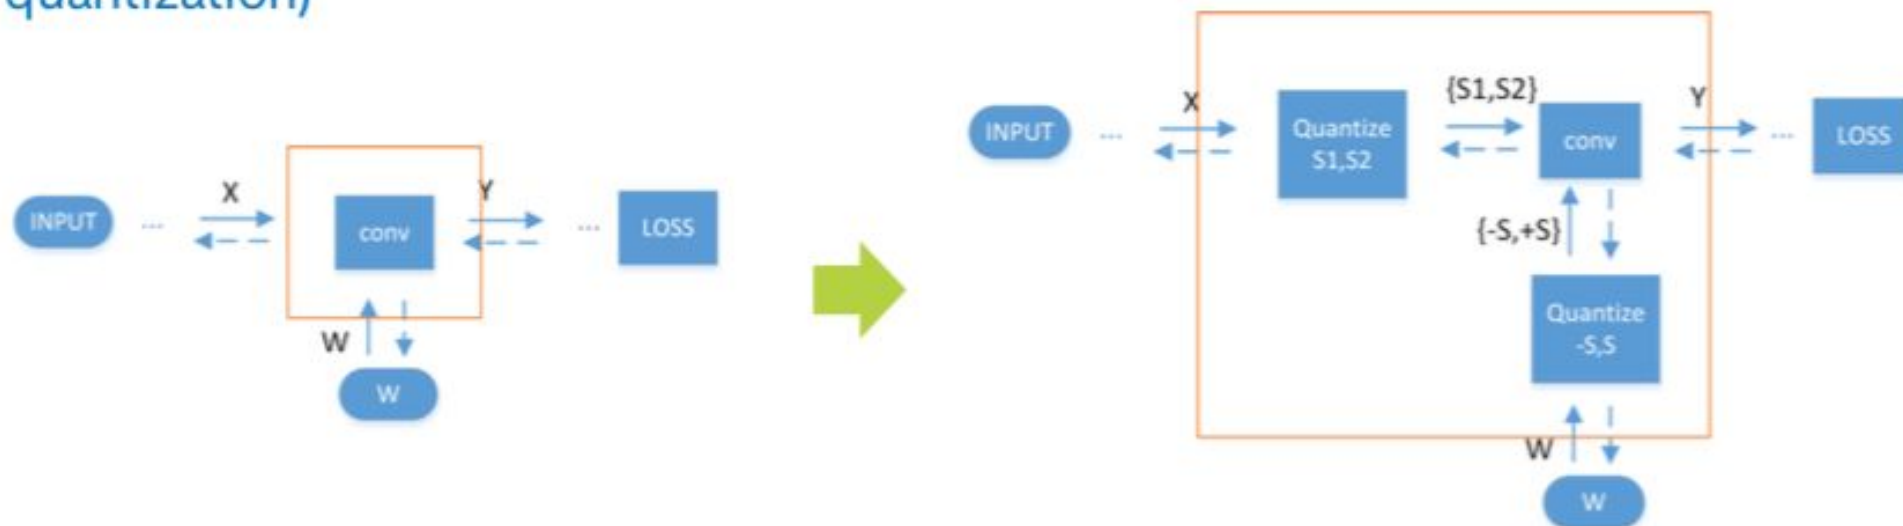
$$a \text{ XNOR } b = \begin{matrix} a \backslash b & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \quad \{0, 1\},$$

$X = [x_1, \dots, 0, 1, 1, \dots, x_{32}]$, $W = [w_1, \dots, w_{32}] \rightarrow \text{int32 number}$
 Calc $y = 2 * \text{popcount}(W \text{ XNOR } X) - 32$
We have only 1 XNOR and 1 POPCOUNT operations


Binary Net

During binarization process selected convolutional layers of the original CNN are replaced with binary convolution alternatives.

Network is modified by inserting special quantization layer for input activations and weights that converts any full-precision value into two pretrained values (so-called “fake” quantization)



Binary Net



Algorithms for Efficient Inference

1. Pruning
2. Weight Sharing
3. Quantization
4. Binary / Ternary Net
- 5. Distillation**
6. Low Rank Approximation
7. Winograd Transformation

Distillation

- **L2 – Ba [14]**
 - L2 loss between teacher and student logits
 - No labels required




Figure 6. Teacher-Student model, L2

- [14] J. Ba and R. Caruana, "Do deep nets really need to be deep?" In Advances in neural information processing systems, 2014, pp. 2654–2662.
- [15] G. Hinton, O. Vinyals, and J. Dean, "Distilling the knowledge in a neural network," ArXiv preprint arXiv:1503.02531, 2015.
- [16] A. Romero, N. Ballas, S. E. Kahou, A. Chassang, C. Gatta, and Y. Bengio, "Fitnets: Hints for thin deep nets," ArXiv preprint arXiv:1412.6550, 2015.

Distillation

- L2 – Ba [14]
 - L2 loss between teacher and student logits
 - No labels required
- **Knowledge Distillation [15]**
 - Soft target: softmax cross entropy with teacher logits
 - Hard target: softmax cross entropy with correct labels




Figure 7. Teacher-Student model, KD

[14] J. Ba and R. Caruana, "Do deep nets really need to be deep?" In Advances in neural information processing systems, 2014, pp. 2654–2662.

[15] G. Hinton, O. Vinyals, and J. Dean, "Distilling the knowledge in a neural network," ArXiv preprint arXiv:1503.02531, 2015.

[16] A. Romero, N. Ballas, S. E. Kahou, A. Chassang, C. Gatta, and Y. Bengio, "Fitnets: Hints for thin deep nets," ArXiv preprint arXiv:1412.6550, 2015.

Distillation

- L2 – Ba [14]
 - L2 loss between teacher and student logits
 - No labels required
- Knowledge Distillation [15]
 - Soft target: softmax cross entropy with teacher logits
 - Hard target: softmax cross entropy with correct labels
- **FitNets [16]**
 - Knowledge Distillation with hints in the middle points of the network
 - Student is deeper than the teacher




Figure 8. Teacher-Student model, FitNets

[14] J. Ba and R. Caruana, "Do deep nets really need to be deep?" In Advances in neural information processing systems, 2014, pp. 2654–2662.

[15] G. Hinton, O. Vinyals, and J. Dean, "Distilling the knowledge in a neural network," ArXiv preprint arXiv:1503.02531, 2015.

[16] A. Romero, N. Ballas, S. E. Kahou, A. Chassang, C. Gatta, and Y. Bengio, "Fitnets: Hints for thin deep nets," ArXiv preprint arXiv:1412.6550, 2015.

Algorithms for Efficient Inference

1. Pruning
2. Weight Sharing
3. Quantization
4. Binary / Ternary Net
5. Distillation
- 6. Low Rank Approximation**
7. Winograd Transformation

Low Rank Approximation

- Basis filter set => Basis feature maps
- Final feature map = linear combination of basis feature maps
- Rank-1 basis filter => decomposed into a sequence of horizontal and vertical filters
- ~2.4x speedup, no performance drop







Figure 1: (a) The original convolutional layer acting on a single-channel input *i.e.* $C=1$. (b) The approximation to that layer using the method of Scheme 1. (c) The approximation to that layer using the method of Scheme 2. Individual filter dimensions are given above the filter layers.



Algorithms for Efficient Inference

1. Pruning
2. Weight Sharing
3. Quantization
4. Binary / Ternary Net
5. Distillation
6. Low Rank Approximation
7. Winograd Transformation

Winograd Transformation



9xC FMAs/Output: Math Intensive

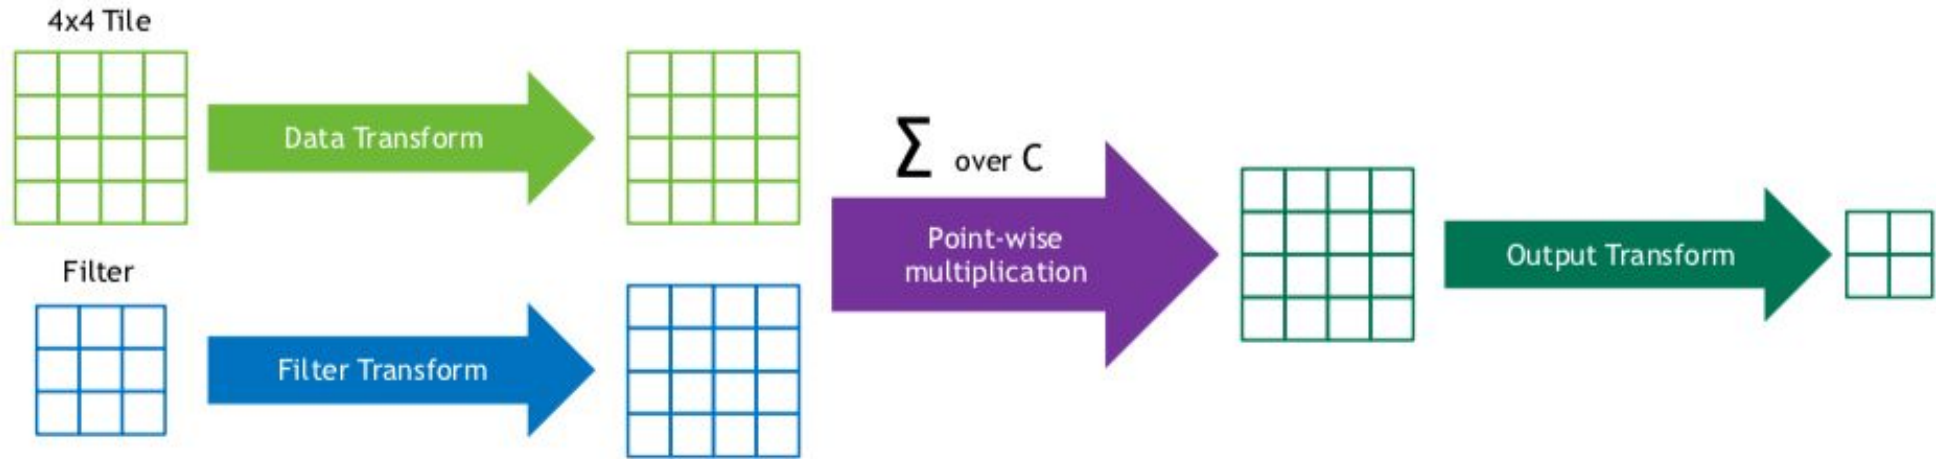
$$\sum \begin{matrix} 0 \times 8 & + & 1 \times 7 & + & 2 \times 6 & + \\ 3 \times 5 & + & 4 \times 4 & + & 5 \times 3 & + \\ 6 \times 2 & + & 7 \times 1 & + & 8 \times 0 \end{matrix}$$

9xK FMAs/Input: Good Data Reuse

4 × 0	4 × 1	4 × 2
4 × 3	4 × 4	4 × 5
4 × 6	4 × 7	4 × 8

Direct convolution: we need $9 \times C \times 4 = 36 \times C$ FMAs for 4 outputs

Winograd Transformation



Direct convolution: we need $9 \times C \times 4 = 36 \times C$ FMAs for 4 outputs

Winograd convolution: we need $16 \times C$ FMAs for 4 outputs: **2.25x** fewer FMAs

Summary

Method	Advantages	Disadvantages
Binarization & Quantization	Low latency and memory usage	High loss of accuracy
Pruning	Prevents overfitting, the accuracy can increase	Slow
Factorization	Can achieve state of the art results while decreasing the computation cost	Dependent on framework
Distillation	Applicable to all architectures Doesn't change the network	only applicable to classification task

Thank you for your
attention!